

# Assignment 1

## MyMyUNSW Schema

Last updated: **Sunday 27th February 3:07pm**

Most recent changes are shown in **red** ... older changes are shown in **brown**.

[\[Assignment Spec\]](#) [\[Database Design\]](#) [\[Schema in SQL\]](#) [\[check1.sql\]](#) [\[Examples\]](#) [\[Fixes+Updates\]](#)

schema.sql

```
-- COMP3311 22T1 Assignment 1 Schema
--
-- MyMyUNSW Schema
-- Original version: John Shepherd (Sept 2002)
-- Latest version: John Shepherd (August 2013)
-- Originally for PostgreSQL7.2
-- Conformed to SQL standard and ported to Oracle8: April 2003
-- Minor mods: April 2005, Sept 2006, April 2007, May 2008
-- Added extra Rule structures: March 2011
-- Adjusted some tables to fit MAPPS/AIMS data structures: August 2013
-- Simplified schema (e.g some tables removed): October 2020
--
-- Gives a standard SQL description for data to maintain information
-- about academic matters at UNSW. Options for simplifying the schema
-- by exploiting non-standard PostgreSQL features are marked with "PG:"
--
-- The notion is that this data should enable all of the functionality
-- currently provided by NSS, CATS, UNSW Staff directory, ...
--
-- To keep the schema a little shorter, I have ignored my usual
-- convention of putting foreign key definitions at the end of
-- the table definition.
--
-- Some general naming principles:
--   max 10 chars in field names
--   all entity tables are named using plural nouns
--   for tables with unique numeric identifier, always call the field "id"
--   for cases where there's a long name and a short name for something,
--       use "name" for the short version of the name (typically for display),
--       and use "longname" for the complete version of the name (which might
--       typically be used in lists of items)
--   for foreign keys referring to an "id" field in the foreign relation,
--       use the singular-noun name of the relation as the field name
--
--       OR use the name of the relationship being represented
--
-- Null values:
--   for each relation, a collection of fields is identified as being
--   compulsory (i.e. without them the data isn't really usable) and
--   they are all defined as NOT NULL
--   reminder: all of the primary keys (e.g. "id") are non-NULL
--   note also that fields that are allowed to be NULL will need to be
--   handled specially whenever they are displayed e.g. in a web-based
--   interface to this schema
```

```
--
-- Enum relations:
--   some relations in the schema contain little more than (id,name)
--   they were not done simply as varchar attributes:
--     for consistency (all relations referring them get common spelling, etc)
--     for efficiency (saves space in the referring relation)
--     for easier use in menus in the user interface
--   examples of such relations: Countries, Room_types, Job_classes, ...
--   you could argue that these should be replaced by PostgreSQL "enum"
--     types, but (a) enums are non-standard, and (b) if you want more
--     info than just a label (e.g. also want a description), you need
--     a table with extra fields
--
-- Meta information:
--   in a couple of cases, the data stored in the database needs to be
--   further interpreted before the actual results can be obtained
-- Examples:
--   Student groups: an SQL query is stored in the DB to extract a
--     list of students in the group
--   Rules: an expression in a simple "requirements language"
--     is stored in the DB and needs to be interpreted by a PLpgSQL
--     function to determine whether the requirements are met
--
-- Oracle port:
--   this schema was converted to standard SQL to run on Oracle in 2003
--   the PostgreSQL non-standard features have been retained as comments

-- Domains: specific kinds of values used throughout
--   In PostgreSQL, some could be defined as simple enumerated types
--   Since we're trying to be standard SQL, we use domains

-- ShortStrings are typically used for values appearing in tables in the UI
create domain ShortString as varchar(16);
create domain MediumString as varchar(64);
create domain LongString as varchar(256);
create domain TextString as varchar(16384);

-- ShortNames are typically used for values appearing in tables in the UI
create domain ShortName as varchar(16);
create domain MediumName as varchar(64);
create domain LongName as varchar(128);

-- If we could rely on having regexps, we could do a better job with these
create domain PhoneNumber as varchar(32);

create domain EmailString as varchar(64) check (value like '%@%');
create domain URLString as varchar(128) check (value like 'http://%');

create domain CareerType as char(2)
    check (value in ('UG','PG','HY','RS','NA'));

create domain GradeType as char(2)
    check (value in (
        'AF', 'AS', 'CR', 'DF', 'DN', 'EC', 'FL', 'FN',
        'GP', 'HD', 'LE', 'NA', 'NC', 'NF', 'PC', 'PE',
```

```
'PS', 'PT', 'RC', 'RD', 'RS', 'SS', 'SY', 'UF',
'WA', 'WC', 'WD', 'WJ', 'XE',
'A', 'B', 'C', 'D', 'E'
));

create domain CampusType as char(1)
    check (value in (
        'K', -- Kensington
        'P', -- COFA/Paddington
        'Z', -- ADFA/UniCollege
        'C', -- CBD (Sydney)
        'X' -- External
    ));

create domain CourseYearType as integer
    check (value > 1945); -- UNSW didn't exist before 1945

create domain TermType as char(2)
    check (value in ('S1','S2','X1','X2','T0','T1','T2','T3'));

-- Countries: country codes and names

create table Countries (
    id            integer, -- PG: serial
    code          char(3) not null unique,
    name          LongName not null,
    primary key (id)
);

-- Buildings: building information
-- e.g. (1234, 'MB', 'Morven Brown Building', 'K', 'C20')
--      (5678, 'K17', 'CSE Building', 'K', 'K17')
--      (4321, 'EE', 'Electrical Engineering Building', 'K', 'G17')

create table Buildings (
    id            integer, -- PG: serial
    unswid        ShortString not null unique,
    name          LongName not null,
    campus        CampusType,
    gridref       char(4),
    primary key (id)
);

-- Room_types: different kinds of rooms on campus
-- e.g. 'Lecture Theatre', 'Tutorial Room', 'Office', ...

create table Room_types (
    id            integer, -- PG: serial
    description MediumString not null,
    primary key (id)
);
```

```
-- Rooms: room information

create table Rooms (
    id            integer, -- PG: serial
    unswid        ShortString not null unique,
    rtype         integer references Room_types(id),
    name          ShortName not null,
    longname      LongName,
    building      integer references Buildings(id),
    capacity      integer check (capacity >= 0),
    primary key (id)
);

-- Facilities: things in rooms (e.g. data projector, OHP, etc.)

create table Facilities (
    id            integer, -- PG: serial
    description   MediumString not null,
    primary key (id)
);

-- Room_facilities: which facilities are available in which rooms

create table Room_facilities (
    room          integer references Rooms(id),
    facility       integer references Facilities(id),
    primary key (room,facility)
);

-- OrgUnit_types: kinds of organisational units at UNSW
-- notes:
--   examples: 'Faculty', 'School', 'Division',...
--   used so that people can invent other new units in the future

create table OrgUnit_types (
    id            integer, -- PG: serial
    name          ShortName not null,
    primary key (id)
);

-- OrgUnits: organisational units (e.g. schools, faculties, ...)
-- notes:
--   "utype" classifies the organisational unit
--

create table OrgUnits (
    id            integer, -- PG: serial
    utype         integer not null references OrgUnit_types(id),
```

```
name      MediumString not null,
longname   LongString,
unswid     ShortString,
phone      PhoneNumber,
email      EmailString,
website    URLString,
starting   date, -- not null
ending     date,
primary key (id)
);

-- OrgUnit_groups: how organisational units are related
-- notes:
--   allows for a multi-level hierarchy of groups

create table OrgUnit_groups (
    owner      integer references OrgUnits(id),
    member     integer references OrgUnits(id),
    primary key (owner,member)
);

-- Teaching Periods (aka terms, sessions, semesters)
-- notes:
--   all dates should be not null, but we don't have access to them

create table Terms (
    id          integer, -- PG: serial
    unswid      integer not null unique,
    year        CourseYearType,
    session     char(2) not null, -- has constraint in database
    name        ShortName not null,
    longname     LongName not null,
    starting    date not null,
    ending      date not null,
    startBrk    date, -- start of mid-semester break
    endBrk      date, -- end of mid-semester break
    endWD       date, -- last date to withdraw without academic penalty
    endEnrol    date, -- last date to enrol without special permission
    census      date, -- last date to withdraw without paying for course
    primary key (id)
);

-- Public_holidays: days when regular teaching is cancelled
-- These could be done as WholeDay/OneOff Events, but they would also
--   need to generate exceptions for all of the Class Events scheduled
--   on those days
-- Notice that there's no primary key; there could be several holidays
--   (e.g. different religions) on the same date

create table Public_holidays (
    term        integer references Terms(id),
```

```

        description MediumString, -- e.g. Good Friday, Easter Day
        day          date
    );

-- Staff_roles: roles for staff within the UNSW organisation
-- handles job classes under which staff are employed
-- e.g. "Associate Lecturer", "Professor", "Administrative Assistant",
--      "Computer Systems Officer", "Clerk", "Caterer"
-- and also handles specific roles for some staff members
-- e.g. "Vice Chancellor", "Dean", "Head of School",
--      "Teaching Director", "Admin Assistant to Dean",
--      "School Office Manager", ...
-- this could either describe the specific duties under the
-- job classification, or duties that are additional to the
-- basic job classification
-- notes:
-- in the real NSS, hooks to the HR system would be here
-- for example, we might have base salary for each role
-- which represent a job classification

create table Staff_role_types (
    id          char(1),
    description ShortString,
    primary key (id)
);

create table Staff_role_classes (
    id          char(1),
    description ShortString,
    primary key (id)
);

create table Staff_roles (
    id          integer, -- PG: serial
    rtype       char(1) references Staff_role_types(id),
    rclass      char(1) references Staff_role_classes(id),
    name        LongString not null,
    description LongString,
    primary key (id)
);

-- People super-class
-- contains:
-- unique id internal to database
-- personal information
-- home contact info
-- notes:
-- family,given names are displayed on transcripts
-- sortname is to handle unusual names (e.g. de Kleer as K)
-- name is what will be displayed (except on transcripts)
-- it allows preferred form of name(s) to be used
-- phone numbers are assumed to be Australian numbers

```

```

-- the phone field sizes allow for future expansion of phone #s
-- familyname is allowed to be null for people with only one name
-- the "not null" fields indicate which info is compulsory
-- nowadays, people are required to have an email address
-- the password field is used by the web interface
-- allows people in the database who are not staff or students
-- e.g. members of the University Council

create table People (
    id            integer, -- PG: serial
    unswid        integer unique, -- staff/student id (can be null)
    password      ShortString not null,
    family        LongName,
    given         LongName not null,
    title         ShortName, -- e.g. "Prof", "A/Prof", "Dr", ...
    sortname      LongName not null,
    name          LongName not null,
    street        LongString,
    city          MediumString,
    state         MediumString,
    postcode      ShortString,
    country       integer references Countries(id),
    homephone     PhoneNumber, -- should be not null
    mobphone      PhoneNumber,
    email         EmailString not null,
    homepage      URLString,
    gender        char(1) check (gender in ('m','f')),
    birthday      date,
    origin        integer references Countries(id), -- country where born
    primary key (id)
);

-- Student (sub-class): enrolment type

create table Students (
    id            integer references People(id),
    stype         varchar(5) check (stype in ('local','intl')),
    primary key (id)
);

-- Student_groups: groups of students (used in specifying quotas)
-- uses SQL queries stored in the database to extract lists of
-- students belonging to particular classes
-- decided to use this approach rather than explicitly storing
-- lists of (student,group) pairs because these lists would
-- be very large and hard to setup and maintain
-- of course, with this approach, getting a list of students
-- in a given group requires something beyond SQL (e.g. PLpgSQL)

create table Student_groups (
    id            integer, -- PG: serial
    name          LongName unique not null,
    definition     TextString not null, -- SQL query to get student(id)'s

```

```
        primary key (id)
    );

-- Staff (sub-class): employment and on-campus contact info
-- all staff have a unique staff id different to their person id
-- anyone who teaches a class has to be entered in this table
-- (they would normally be entered into the UNSW HR database)

create table Staff (
    id            integer references People(id),
    office        integer references Rooms(id),
    phone         PhoneNumber, -- full number, not just extension
    employed      date not null,
    supervisor    integer references Staff(id),
    primary key (id)
);

-- Affiliations: staff roles and association to organisational units
-- notes:
-- most staff will be attached to only one unit
-- "role" will describe things like "Professor", "Head of School", ...
-- if this is their job class for HR, isPrimary is true

create table Affiliations (
    staff         integer references Staff(id),
    orgUnit       integer references OrgUnits(id),
    role          integer references Staff_roles(id),
    isPrimary      boolean, -- is this role the basis for their employment?
    starting      date not null, -- when they commenced this role
    ending        date, -- when they finished; null means current
    primary key (staff,orgUnit,role,starting)
);

-- Programs: academic details of a degree program
-- notes:
-- the "code" field is used for compatability with current UNSW practice
-- e.g. 3978 is the code for the computer science degree

create table Programs (
    id            integer, -- PG: serial
    code          char(4) not null, -- e.g. 3978, 3645, 3648
    name          LongName not null,
    uoc           integer check (uoc >= 0),
    offeredBy     integer references OrgUnits(id),
    career        CareerType,
    duration       integer, -- #months
    description    TextString, -- PG: text
    firstOffer    integer references Terms(id), -- should be not null
    lastOffer     integer references Terms(id), -- null means current
    primary key (id)
);
```



```
-- Streams: academic details of a major/minor stream(s) in a degree

create table Streams (
    id            integer, -- PG: serial
    code          char(6) not null, -- e.g. COMPA1, SENGAI
    name          LongName not null,
    offeredBy     integer references OrgUnits(id),
    stype         ShortString,
    description    TextString,
    firstOffer    integer references Terms(id), -- should be not null
    lastOffer     integer references Terms(id), -- null means current
    primary key (id)
);

-- Degree_types: types of awards for degrees

create table Degree_types (
    id            integer, -- PG: serial
    unswid        ShortName not null unique, -- e.g. BSc, BSc(CompSci), BE, PhD
    name          MediumString not null, -- e.g. Bachelor of Science
    prefix        MediumString,
    career        CareerType,
    aqf_level     integer check (aqf_level > 0),
    primary key (id)
);

-- Program_degrees: degrees awarded for each program
--   a concurrent degree will have two entries for one program

create table Program_degrees (
    program       integer references Programs(id),
    degree        integer references Degree_types(id),
    name          LongString not null,
    abbrev        MediumString,
    primary key (program,degree)
);

-- Degrees_awarded: info about student being awarded a degree

create table Degrees_awarded (
    student       integer references Students(id),
    program       integer references Programs(id),
    graduated     date,
    primary key (student,program)
);

-- Academic_standing: kinds of academic standing at UNSW
-- e.g. 'good', 'probation1', 'probation2',...
```

```
-- An enumerated-type table

create table Academic_standing (
    id            integer,
    standing      ShortName not null,
    notes         TextString,
    primary key (id)
);

-- Subjects: academic details of a course (version)
-- "code" is standard UNSW course code (e.g. COMP3311)
-- "firstOffer" and "lastOffer" indicate a timespan during
--   which this subject was offered to students; if "lastOffer"
--   is null, then the subject is still running
-- Note: UNSW calls subjects "courses"

create table Subjects (
    id            integer, -- PG: serial
    code          char(8) not null,
    --           PG: check (code ~ '[A-Z]{4}[0-9]{4}'),
    name          MediumName not null,
    longname      LongName,
    uoc           integer check (uoc >= 0),
    offeredBy     integer references OrgUnits(id),
    eftsload      float,
    career        CareerType,
    syllabus      TextString, -- PG: text
    contactHPW    float, -- contact hours per week
    _excluded     text,    -- plain text from MAPPS
    excluded      integer, -- references Acad_object_groups(id),
    _equivalent   text,    -- plain text from MAPPS
    equivalent    integer, -- references Acad_object_groups(id),
    _prereq       text,    -- plain text from MAPPS
    prereq        integer, -- references Rules(id)
    replaces      integer references Subjects(id),
    firstOffer    integer references Terms(id), -- should be not null
    lastOffer     integer references Terms(id), -- null means current
    primary key (id)
);

-- Course: info about an offering of a subject in a given term
-- we insist on knowing the lecturer because there's no point running
--   a course unless you've got someone organised to lecture it
-- Note: UNSW calls courses "course offerings"

create table Courses (
    id            integer, -- PG: serial
    subject       integer not null references Subjects(id),
    term          integer not null references Terms(id),
    homepage      URLString,
    primary key (id)
);
```

```
-- Course_staff: various staff involved in a course
-- allows one Staff to have multiple roles in a course

create table Course_staff (
    course      integer references Courses(id),
    staff       integer references Staff(id),
    role        integer references Staff_roles(id),
    primary key (course,staff,role)
);

-- Course_quotas: quotas for various classes of students in a course
-- if there's no quota, there's no entry in this table
-- alternatively, we could have allowed quota to be null
-- and used that as a mechanism for indicating "no quota"

create table Course_quotas (
    course      integer references Courses(id),
    sgroup      integer references Student_groups(id),
    quota       integer not null,
    primary key (course,sgroup)
);

-- Program_enrolments: student's enrolment in a program in one term
-- notes:
--   "standing" refers to the students academic standing
--   "wam" is computed from marks in enrolment records

create table Program_enrolments (
    id          integer,
    student     integer not null references Students(id),
    term        integer not null references Terms(id),
    program     integer not null references Programs(id),
    wam         real,
    standing    integer references Academic_standing(id),
    advisor     integer references Staff(id),
    notes       TextString,
    primary key (id)
);

-- Stream_enrolments: student's enrolment in streams in one term

create table Stream_enrolments (
    partOf      integer references Program_enrolments(id),
    stream      integer references Streams(id),
    primary key (partOf,stream)
);

-- Course_enrolments: student's enrolment in a course offering
```

```
-- null grade means "currently enrolled"
-- if course is graded SY/FL, then mark always remains null

create table Course_enrolments (
    student    integer references Students(id),
    course     integer references Courses(id),
    mark       integer check (mark >= 0 and mark <= 100),
    grade      GradeType,
    stuEval    integer check (stuEval >= 1 and stuEval <= 6),
    primary key (student,course)
);

-- Books: textbook details

create table Books (
    id         integer, -- PG: serial
    isbn       varchar(20) unique,
    title      LongString not null,
    authors    LongString not null,
    publisher  LongString not null,
    edition    integer,
    pubYear    integer not null check (pubYear > 1900),
    primary key (id)
);

-- Course_books: relates books to courses
-- books are related to a Course rather than a Subject because texts
-- may change over time, even if the syllabus remains constant

create table Course_books (
    course     integer references Courses(id),
    book       integer references Books(id),
    bktype     varchar(10) not null check (bktype in ('Text','Reference')),
    primary key (course,book)
);

-- ClassType: names for different kinds of class
-- e.g. "Lecture", "Tutorial", "Lab Class", ...

create table Class_types (
    id         integer, -- PG: serial
    unswid     ShortString not null unique,
    name       MediumName not null,
    description MediumString,
    primary key (id)
);

-- Classes: a specific regular teaching event in a course
-- we ignore streams, since they make class registration too messy
-- we don't allow day/time/place info to be null; this forces us to
```

```
-- already organise a time/place before we enter them in the system
-- weekly repetitions are handled by (repeats=1 or repeats is null)
-- we assume that all classes are multiples of 1-hour in duration
-- and cannot start before 8am or finish after 11pm)

create table Classes (
    id            integer, -- PG: serial
    course        integer not null references Courses(id),
    room          integer not null references Rooms(id),
    ctype         integer not null references Class_types(id),
    dayOfWk       integer not null check (dayOfWk >= 0 and dayOfWk <= 6),
                                -- Sun=0 Mon=1 Tue=2 ... Sat=6
    startTime     integer not null check (startTime >= 8 and startTime <= 22),
    endTime       integer not null check (endTime >= 9 and endTime <= 23),
                                -- time of day, between 8am and 11pm
    startDate     date not null,
    endDate       date not null,
    repeats       integer, -- every X weeks
    primary key (id)
);

-- Class_teachers: who teaches which class
-- unfortunately, no way to describe how two staff who
-- are allocated to a given class teach together
-- e.g. teach on alternating weeks

create table Class_teachers (
    class         integer references Classes(id),
    teacher       integer references Staff(id),
    primary key (class,teacher)
);

-- Class_enrolments: one student's enrolment in a class

create table Class_enrolments (
    student       integer references Students(id),
    class         integer references Classes(id),
    primary key (student,class)
);

-- External_subjects: represents courses from other institutions
-- used to ensure consistency in awarding advanced standing
-- if student X gets advanced standing based on course Y at Z,
-- then a later student who has done course Y at Z can be given
-- the same advanced standing
-- to do this properly, we'd need to set up a table of external
-- institutions and use a foreign key ... as it stands, if
-- people award credit for the same course, but spell either
-- the course name or the institution name differently, it
-- will be treated as a different course
```

```

create table External_subjects (
    id            integer,
    extsubj       LongName not null,
    institution    LongName not null,
    yearOffered   CourseYearType,
    equivTo       integer not null references Subjects(id),
-- creator       integer not null references Staff(id),
-- created       date not null,
    primary key (id)
);

-- Variations: replacement of one subject or another in a program
-- handles several cases (which are more or less similar):
--   advanced standing for courses studied either at UNSW or elsewhere
--   substitution of one course for another to satisfy requirements
--   exemption from one course, to use as a prerequisite
-- in the case of exemptions, no credit is granted towards a program;
--   the subject is being recorded to use as a pre-req
-- the substitution is for one subject towards the requirements
--   of one stream
-- there are two sub-cases represented in this single table:
--   the subject is an internal UNSW subject (internal equivalence)
--   the subject is from outside UNSW (external equivalence)
-- can't enter Advanced Standing without saying who you are, since
--   Advanced Standing is like awarding a pass in a UNSW course
-- if we wanted to record external subjects being used as a basis
--   for pre-requisites but not credit (i.e. exemption), we would
--   need to add a new field to indicate that no credit was involved

create domain VariationType as ShortName
    check (value in ('advstanding','substitution','exemption'));

create table Variations (
    student       integer references Students(id),
    program       integer references Programs(id),
    subject       integer references Subjects(id),
    vtype         VariationType not null,
    intEquiv      integer references Subjects(id),
    extEquiv      integer references External_subjects(id),
    yearPassed    CourseYearType,
    mark          integer check (mark > 0), -- if we know it
    approver      integer not null references Staff(id),
    approved      date not null,

    primary key (student,program,subject),
    constraint TwoCases check
        ((intEquiv is null and extEquiv is not null)
         or
         (intEquiv is not null and extEquiv is null))
);

-- Acad_object_groups: groups of different kinds of academic objects
-- academic objects = courses OR streams OR programs

```

```
-- different kinds of academic objects that can be grouped
-- each group consists of a set of objects of the same type

create domain AcadObjectGroupType as ShortName
    check (value in (
        'subject',      -- group of subjects
        'stream',       -- group of streams
        'program'       -- group of programs
    ));

-- how to interpret combinations of objects in groups

create domain AcadObjectGroupLogicType as ShortName
    check (value in ( 'and', 'or' ));

-- how groups are defined

create domain AcadObjectGroupDefType as ShortName
    check (value in ( 'enumerated', 'pattern', 'query' ));

-- there are some constraints in this table that we haven't implemented

create table Acad_object_groups (
    id          integer,
    name        LongName,
    gtype       AcadObjectGroupType not null,
    glogic       AcadObjectGroupLogicType,
    gdefBy       AcadObjectGroupDefType not null,
    negated      boolean default false,
    parent       integer, -- references Acad_object_groups(id),
    definition   TextString, -- if pattern or query-based group
    primary key (id)
);

alter table Acad_object_groups
    add foreign key (parent) references Acad_object_groups(id);

alter table Subjects
    add foreign key (excluded) references Acad_object_groups(id);

alter table Subjects
    add foreign key (equivalent) references Acad_object_groups(id);

-- Each kind of AcademicObjectGroup requires it own membership relation

create table Subject_group_members (
    subject      integer references Subjects(id),
    ao_group     integer references Acad_object_groups(id),
    primary key (subject,ao_group)
);

create table Stream_group_members (
    stream       integer references Streams(id),
    ao_group     integer references Acad_object_groups(id),
```

```

        primary key (stream,ao_group)
    );

create table Program_group_members (
    program      integer references Programs(id),
    ao_group     integer references Acad_object_groups(id),
    primary key (program,ao_group)
);

-- Rules: requirements for programs and stream, pre-reqs for subjects

create domain RuleType as char(2)
    check (value in (
        'CC', -- core courses ... with min, max, subject group
        'PE', -- program electives ... with min, max, subject group
        'FE', -- free electives ... with min, max, group with FREE?###
        'GE', -- general education ... with min, max, group with GEN??###
        'RQ', -- subject pre-req ... typically with min, max, subject group
        'WM', -- WAM requirement ... typically with min WAM score
        'LR', -- limit rule ... with min or max, big subject group (####...)
        'MR', -- maturity rule ... with min UOC and (optionally) a subject group
        'DS', -- done stream ... with min, max, stream group
        'RC', -- recommended ... with subject group, useful for suggestions
        'IR'  -- information rule ... doesn't need checking
    ));

-- Various types of rules ...
-- Some rules require reference to a group of subjects or streams
-- min/max can have different kinds of units depending on rule type
--   (frequently they are UOC, sometimes just counters)
-- Rule names don't have a standard form and are not very useful
-- Rule descriptions are slightly more useful

create table Rules (
    id          integer,
    name        MediumName,
    type        RuleType,
    min         integer check (min >= 0),
    max         integer check (min >= 0),
    ao_group    integer references Acad_object_groups(id),
    description TextString,
    primary key (id)
);

create table Subject_prereqs (
    subject     integer references Subjects(id),
    career      CareerType, -- what kind of students it applies to
    rule        integer references Rules(id),
    primary key (subject,career,rule)
);

create table Stream_rules (
    stream      integer references Streams(id),
    rule        integer references Rules(id),

```



```
        primary key (stream,rule)
    );

create table Program_rules (
    program    integer references Programs(id),
    rule       integer references Rules(id),
    primary key (program,rule)
);
```

